

# Data Transformation

Gabriel Montes

2025-07-17

## Exercise 3.2.5

1. In a single pipeline for each condition, find all flights that meet the condition:

```
library(nycflights13)
library(dplyr)
```

Had an arrival delay of two or more hours

```
##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
flights |>
  filter(arr_delay >= 120)
```

```
## # A tibble: 10,200 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     811             630        101    1047             830
## 2  2013     1     1     848             1835       853    1001             1950
## 3  2013     1     1     957             733        144    1056             853
## 4  2013     1     1    1114             900        134    1447             1222
## 5  2013     1     1    1505             1310       115    1638             1431
## 6  2013     1     1    1525             1340       105    1831             1626
## 7  2013     1     1    1549             1445         64    1912             1656
## 8  2013     1     1    1558             1359       119    1718             1515
## 9  2013     1     1    1732             1630         62    2028             1825
## 10 2013     1     1    1803             1620       103    2008             1750
## # i 10,190 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights |>
  filter(dest == "IAH" | dest == "HOU")
```

## Flew to Houston (IAH or HOU)

```
## # A tibble: 9,313 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     623           627          -4     933           932
## 4  2013     1     1     728           732          -4    1041          1038
## 5  2013     1     1     739           739           0    1104          1038
## 6  2013     1     1     908           908           0    1228          1219
## 7  2013     1     1    1028          1026           2    1350          1339
## 8  2013     1     1    1044          1045          -1    1352          1351
## 9  2013     1     1    1114           900        134    1447          1222
## 10 2013     1     1    1205          1200           5    1503          1505
## # i 9,303 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights |>
  filter(carrier %in% c("UA", "AA", "DL"))
```

## Were operated by United, American, or Delta

```
## # A tibble: 139,504 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     554           600          -6     812           837
## 5  2013     1     1     554           558          -4     740           728
## 6  2013     1     1     558           600          -2     753           745
## 7  2013     1     1     558           600          -2     924           917
## 8  2013     1     1     558           600          -2     923           937
## 9  2013     1     1     559           600          -1     941           910
## 10 2013     1     1     559           600          -1     854           902
## # i 139,494 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights |>
  filter(month %in% c(7, 8, 9))
```

## Departed in summer (July, August, and September)

```
## # A tibble: 86,326 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     7     1         1           2029        212     236          2359
## 2  2013     7     1         2           2359         3     344           344
## 3  2013     7     1        29           2245        104     151           1
```

```
## 4 2013 7 1 43 2130 193 322 14
## 5 2013 7 1 44 2150 174 300 100
## 6 2013 7 1 46 2051 235 304 2358
## 7 2013 7 1 48 2001 287 308 2305
## 8 2013 7 1 58 2155 183 335 43
## 9 2013 7 1 100 2146 194 327 30
## 10 2013 7 1 100 2245 135 337 135
## # i 86,316 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights |>
  filter(arr_delay > 120, dep_delay <= 0)
```

Arrived more than two hours late but didn't leave late

```
## # A tibble: 29 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1 2013     1    27   1419           1420        -1    1754           1550
## 2 2013    10     7   1350           1350         0    1736           1526
## 3 2013    10     7   1357           1359        -2    1858           1654
## 4 2013    10    16    657            700        -3    1258           1056
## 5 2013    11     1    658            700        -2    1329           1015
## 6 2013     3    18   1844           1847        -3         39           2219
## 7 2013     4    17   1635           1640        -5    2049           1845
## 8 2013     4    18    558            600        -2    1149            850
## 9 2013     4    18    655            700        -5    1213            950
## 10 2013     5    22   1827           1830        -3    2217           2010
## # i 19 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights |>
  filter(dep_delay >= 60, dep_delay - arr_delay > 30)
```

Were delayed by at least an hour, but made up over 30 minutes in flight

```
## # A tibble: 1,844 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1 2013     1     1   2205           1720       285         46           2040
## 2 2013     1     1   2326           2130       116        131            18
## 3 2013     1     3   1503           1221       162       1803           1555
## 4 2013     1     3   1839           1700        99       2056           1950
## 5 2013     1     3   1850           1745        65       2148           2120
## 6 2013     1     3   1941           1759       102       2246           2139
## 7 2013     1     3   1950           1845        65       2228           2227
## 8 2013     1     3   2015           1915        60       2135           2111
## 9 2013     1     3   2257           2000       177         45           2224
## 10 2013     1     4   1917           1700       137       2135           1950
```

```
## # i 1,834 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

2. Sort flights to find the flights with the longest departure delays. Find the flights that left earliest in the morning.

```
flights |>
  arrange(desc(dep_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     9     641             900         1301    1242         1530
## 2  2013     6    15    1432            1935         1137    1607         2120
## 3  2013     1    10    1121            1635         1126    1239         1810
## 4  2013     9    20    1139            1845         1014    1457         2210
## 5  2013     7    22     845            1600         1005    1044         1815
## 6  2013     4    10    1100            1900          960    1342         2211
## 7  2013     3    17    2321             810          911     135         1020
## 8  2013     6    27     959            1900          899    1236         2226
## 9  2013     7    22    2257             759          898     121         1026
## 10 2013    12     5     756            1700          896    1058         2020
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

3. Sort flights to find the fastest flights. (Hint: Try including a math calculation inside of your function.)

```
flights |>
  mutate(speed = distance / (air_time / 60)) |>
  arrange(desc(speed))
```

```
## # A tibble: 336,776 x 20
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     5    25    1709             1700          9    1923         1937
## 2  2013     7     2    1558             1513          45    1745         1719
## 3  2013     5    13    2040             2025          15    2225         2226
## 4  2013     3    23    1914             1910           4    2045         2043
## 5  2013     1    12    1559             1600          -1    1849         1917
## 6  2013    11    17     650             655          -5    1059         1150
## 7  2013     2    21    2355             2358          -3     412          438
## 8  2013    11    17     759             800          -1    1212         1255
## 9  2013    11    16    2003             1925          38     17           36
## 10 2013    11    16    2349             2359         -10     402          440
## # i 336,766 more rows
## # i 12 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>, speed <dbl>
```

#### 4. Was there a flight on every day of 2013?

```
flights |>
  count(year, month, day) |>
  nrow() == 365
```

```
## [1] TRUE
```

#### 5. Which flights traveled the farthest distance? Which traveled the least distance?

```
flights |>
  arrange(desc(distance))
```

##### Farthest flight

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     857             900          -3    1516         1530
## 2  2013     1     2     909             900           9    1525         1530
## 3  2013     1     3     914             900          14    1504         1530
## 4  2013     1     4     900             900           0    1516         1530
## 5  2013     1     5     858             900          -2    1519         1530
## 6  2013     1     6    1019             900          79    1558         1530
## 7  2013     1     7    1042             900         102    1620         1530
## 8  2013     1     8     901             900           1    1504         1530
## 9  2013     1     9     641             900        1301    1242         1530
##10  2013     1    10     859             900          -1    1449         1530
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights |>
  arrange(distance)
```

##### Shortest flight

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     7    27      NA             106          NA      NA           245
## 2  2013     1     3    2127             2129          -2    2222         2224
## 3  2013     1     4    1240             1200          40    1333         1306
## 4  2013     1     4    1829             1615         134    1937         1721
## 5  2013     1     4    2128             2129          -1    2218         2224
## 6  2013     1     5    1155             1200          -5    1241         1306
## 7  2013     1     6    2125             2129          -4    2224         2224
## 8  2013     1     7    2124             2129          -5    2212         2224
## 9  2013     1     8    2127             2130          -3    2304         2225
##10  2013     1     9    2126             2129          -3    2217         2224
## # i 336,766 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
```

```
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

6. Does it matter what order you used `filter()` and `arrange()` if you're using both? Why/why not? Think about the results and how much work the functions would have to do.

so if you use `filter()` first it filters the dataset thus making it smaller, if you use `arrange()` first, you are sorting the whole dataset, so order matters

### Exercise 3.3.5

1. Compare `dep_time`, `sched_dep_time`, and `dep_delay`. How would you expect those three numbers to be related?

```
flights |>
  select(dep_time, sched_dep_time, dep_delay) |>
  mutate(
    dep_time_mins = (dep_time %/% 100) * 60 + (dep_time %% 100),
    sched_dep_time_mins = (sched_dep_time %/% 100) * 60 + (sched_dep_time %% 100),
    calculated_delay = dep_time_mins - sched_dep_time_mins
  )
```

```
## # A tibble: 336,776 x 6
##   dep_time sched_dep_time dep_delay dep_time_mins sched_dep_time_mins
##   <int>      <int>      <dbl>      <dbl>          <dbl>
## 1     517        515         2         317            315
## 2     533        529         4         333            329
## 3     542        540         2         342            340
## 4     544        545        -1         344            345
## 5     554        600        -6         354            360
## 6     554        558        -4         354            358
## 7     555        600        -5         355            360
## 8     557        600        -3         357            360
## 9     557        600        -3         357            360
## 10    558        600        -2         358            360
## # i 336,766 more rows
## # i 1 more variable: calculated_delay <dbl>
```

2. Brainstorm as many ways as possible to select `dep_time`, `dep_delay`, `arr_time`, and `arr_delay` from `flights`.

```
flights |> select(dep_time, dep_delay, arr_time, arr_delay)
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>    <dbl>    <int>    <dbl>
## 1     517         2     830        11
## 2     533         4     850        20
## 3     542         2     923        33
## 4     544        -1    1004       -18
## 5     554        -6     812       -25
## 6     554        -4     740        12
## 7     555        -5     913        19
## 8     557        -3     709       -14
## 9     557        -3     838        -8
## 10    558        -2     753         8
```

```
## # i 336,766 more rows
```

```
flights |> select(ends_with("time"), ends_with("delay"))
```

```
## # A tibble: 336,776 x 7
```

```
##   dep_time sched_dep_time arr_time sched_arr_time air_time dep_delay arr_delay
##   <int>      <int>      <int>      <int>      <dbl>    <dbl>    <dbl>
## 1     517         515        830         819        227         2         11
## 2     533         529        850         830        227         4         20
## 3     542         540        923         850        160         2         33
## 4     544         545       1004        1022        183        -1        -18
## 5     554         600        812         837        116        -6        -25
## 6     554         558        740         728        150        -4         12
## 7     555         600        913         854        158        -5         19
## 8     557         600        709         723         53        -3        -14
## 9     557         600        838         846        140        -3         -8
## 10    558         600        753         745        138        -2          8
```

```
## # i 336,766 more rows
```

3. What happens if you specify the name of the same variable multiple times in a select() call?

```
flights |> select(dep_time, dep_time, dep_time)
```

```
## # A tibble: 336,776 x 1
```

```
##   dep_time
##   <int>
## 1     517
## 2     533
## 3     542
## 4     544
## 5     554
## 6     554
## 7     555
## 8     557
## 9     557
## 10    558
```

```
## # i 336,766 more rows
```

they are ignored

4. What does the any\_of() function do? Why might it be helpful in conjunction with this vector?

```
vars <- c("year", "month", "day", "not_a_column")
flights |> select(any_of(vars))
```

```
## # A tibble: 336,776 x 3
```

```
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
```

```
## 7 2013 1 1
## 8 2013 1 1
## 9 2013 1 1
## 10 2013 1 1
## # i 336,766 more rows
```

it is safer to use than `all_of`

5. Does the result of running the following code surprise you? How do the select helpers deal with upper and lower case by default? How can you change that default?

```
flights |> select(contains("TIME"))
```

```
## # A tibble: 336,776 x 6
##   dep_time sched_dep_time arr_time sched_arr_time air_time time_hour
##   <int>         <int>    <int>         <int>      <dbl> <dtm>
## 1     517           515      830           819      227 2013-01-01 05:00:00
## 2     533           529      850           830      227 2013-01-01 05:00:00
## 3     542           540      923           850      160 2013-01-01 05:00:00
## 4     544           545     1004          1022      183 2013-01-01 05:00:00
## 5     554           600      812           837      116 2013-01-01 06:00:00
## 6     554           558      740           728      150 2013-01-01 05:00:00
## 7     555           600      913           854      158 2013-01-01 06:00:00
## 8     557           600      709           723       53 2013-01-01 06:00:00
## 9     557           600      838           846      140 2013-01-01 06:00:00
## 10    558           600      753           745      138 2013-01-01 06:00:00
## # i 336,766 more rows
```

it is case sensitive

```
flights |> select(contains("time", ignore.case = TRUE))
```

```
## # A tibble: 336,776 x 6
##   dep_time sched_dep_time arr_time sched_arr_time air_time time_hour
##   <int>         <int>    <int>         <int>      <dbl> <dtm>
## 1     517           515      830           819      227 2013-01-01 05:00:00
## 2     533           529      850           830      227 2013-01-01 05:00:00
## 3     542           540      923           850      160 2013-01-01 05:00:00
## 4     544           545     1004          1022      183 2013-01-01 05:00:00
## 5     554           600      812           837      116 2013-01-01 06:00:00
## 6     554           558      740           728      150 2013-01-01 05:00:00
## 7     555           600      913           854      158 2013-01-01 06:00:00
## 8     557           600      709           723       53 2013-01-01 06:00:00
## 9     557           600      838           846      140 2013-01-01 06:00:00
## 10    558           600      753           745      138 2013-01-01 06:00:00
## # i 336,766 more rows
```

to make it non sensitive

6. Rename `air_time` to `air_time_min` to indicate units of measurement and move it to the beginning of the data frame.

```
flights |>
  rename(air_time_min = air_time) |>
  relocate(air_time_min)
```



```
## # A tibble: 336,776 x 19
##   air_time_min year month   day dep_time sched_dep_time dep_delay arr_time
##           <dbl> <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1           227  2013     1     1     517           515           2     830
## 2           227  2013     1     1     533           529           4     850
## 3           160  2013     1     1     542           540           2     923
## 4           183  2013     1     1     544           545          -1    1004
## 5           116  2013     1     1     554           600          -6     812
## 6           150  2013     1     1     554           558          -4     740
## 7           158  2013     1     1     555           600          -5     913
## 8            53  2013     1     1     557           600          -3     709
## 9           140  2013     1     1     557           600          -3     838
## 10          138  2013     1     1     558           600          -2     753
## # i 336,766 more rows
## # i 11 more variables: sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

## 7. Why doesn't the following work, and what does the error mean?

```
flights |> select(tailnum) |> arrange(arr_delay)
```

to fix it, because the `arr_delay` is removed, we need to include it in `select`

```
flights |>
  select(tailnum, arr_delay) |>
  arrange(arr_delay)
```

```
## # A tibble: 336,776 x 2
##   tailnum arr_delay
##   <chr>         <dbl>
## 1 N843VA         -86
## 2 N840VA         -79
## 3 N851UA         -75
## 4 N3KCAA         -75
## 5 N551AS         -74
## 6 N24212         -73
## 7 N3760C         -71
## 8 N806UA         -71
## 9 N805JB         -71
## 10 N855VA        -70
## # i 336,766 more rows
```

## Exercise 3.5.7

1. Which carrier has the worst average delays? Challenge: can you disentangle the effects of bad airports vs. bad carriers? Why/why not? (Hint: think about `flights |> group_by(carrier, dest) |> summarize(n())`)

```
library(nycflights13)
library(dplyr)

flights |>
  group_by(carrier) |>
  summarize(avg_delay = mean(dep_delay, na.rm = TRUE)) |>
```

```
arrange(desc(avg_delay)) |>
head()
```

```
## # A tibble: 6 x 2
##   carrier avg_delay
##   <chr>      <dbl>
## 1 F9        20.2
## 2 EV        20.0
## 3 YV        19.0
## 4 FL        18.7
## 5 WN        17.7
## 6 9E        16.7
```

to untangle

```
flights |>
  group_by(carrier, dest) |>
  summarize(avg_delay = mean(dep_delay, na.rm = TRUE), n = n()) |>
  filter(n > 50) |>
  arrange(desc(avg_delay))
```

```
## `summarise()` has grouped output by 'carrier'. You can override using the
## `.groups` argument.
```

```
## # A tibble: 254 x 4
## # Groups:   carrier [15]
##   carrier dest avg_delay    n
##   <chr>   <chr>   <dbl> <int>
## 1 EV      TYS      41.8   323
## 2 EV      CAE      36.7   113
## 3 EV      TUL      34.9   315
## 4 WN      MSY      33.4   298
## 5 EV      OKC      30.6   346
## 6 EV      BHM      29.7   297
## 7 EV      DSM      28.8   478
## 8 EV      TVC      27.5    68
## 9 9E      CLE      25.6   349
## 10 EV     PWM      25.3   813
## # i 244 more rows
```

## 2. Find the flights that are most delayed upon departure from each destination.

```
flights |>
  group_by(dest) |>
  slice_max(dep_delay, n = 1, with_ties = FALSE) |>
  select(dest, carrier, flight, dep_delay)
```

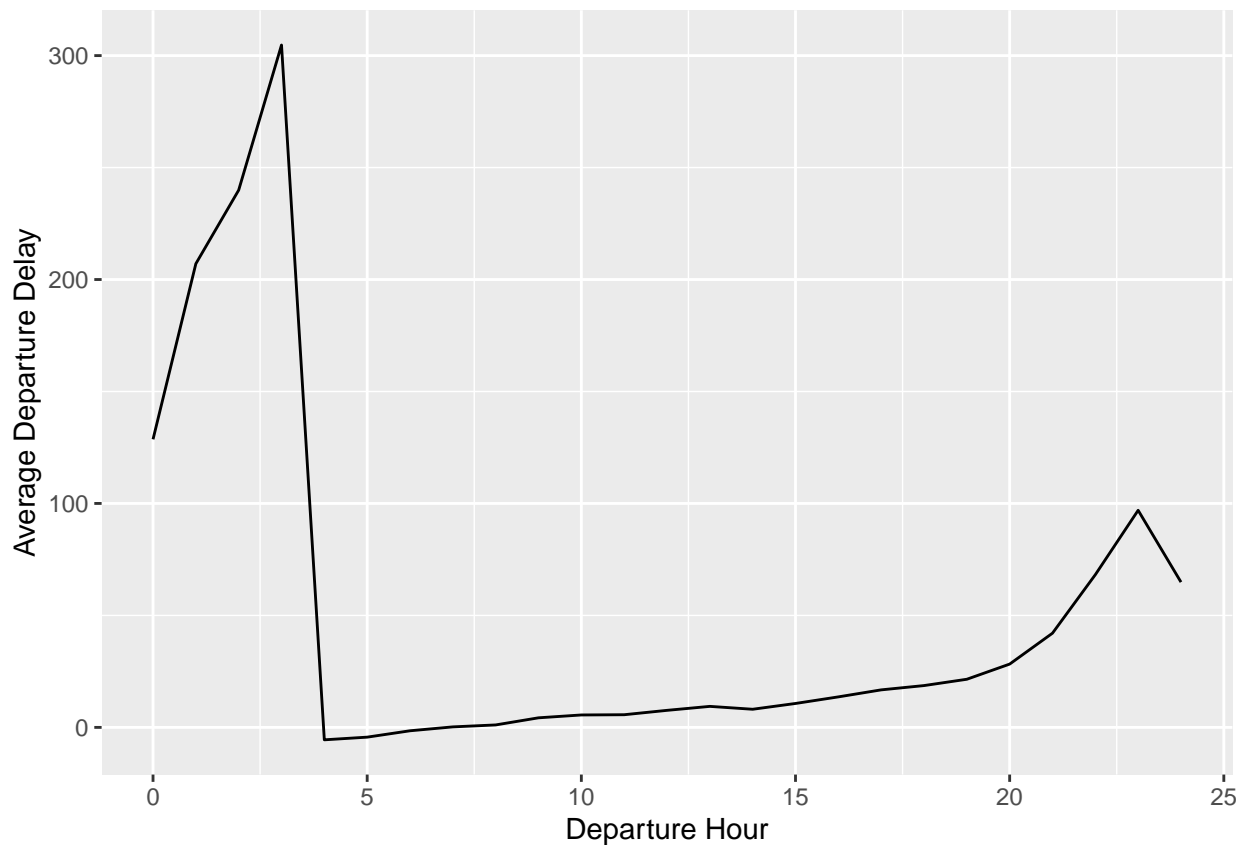
```
## # A tibble: 105 x 4
## # Groups:   dest [105]
##   dest carrier flight dep_delay
##   <chr> <chr>   <int>   <dbl>
## 1 ABQ   B6        65     142
## 2 ACK   B6     1491     219
## 3 ALB   EV     4309     323
## 4 ANC   UA      887      75
## 5 ATL   DL     2047     898
```

```
## 6 AUS UA 503 351
## 7 AVL EV 4519 222
## 8 BDL EV 4103 252
## 9 BGR EV 5309 248
## 10 BHM EV 5038 325
## # i 95 more rows
```

3. How do delays vary over the course of the day? Illustrate your answer with a plot.

```
library(ggplot2)
flights |>
  mutate(hour = dep_time %/% 100) |>
  group_by(hour) |>
  summarize(avg_delay = mean(dep_delay, na.rm = TRUE),
            count = n()) |>
  ggplot(aes(x = hour, y = avg_delay)) +
    geom_line() +
    labs(x = "Departure Hour", y = "Average Departure Delay")
```

```
## Warning: Removed 1 row containing missing values or values outside the scale range
## (`geom_line()`).
```



4. What happens if you supply a negative n to slice\_min() and friends?

they will be ignored as they use absolute values

5. Explain what `count()` does in terms of the `dplyr` verbs you just learned. What does the `sort` argument to `count()` do?

`count()` is like grouping and then sorting is sorting, so you get the most common characteristics

6. Suppose we have the following tiny data frame:

```
df <- tibble(  
  x = 1:5,  
  y = c("a", "b", "a", "a", "b"),  
  z = c("K", "K", "L", "L", "K")  
)
```

a. Write down what you think the output will look like, then check if you were correct, and describe what `group_by()` does.

```
df |>  
  group_by(y)  
  
## # A tibble: 5 x 3  
## # Groups:   y [2]  
##       x y     z  
##   <int> <chr> <chr>  
## 1     1 a     K  
## 2     2 b     K  
## 3     3 a     L  
## 4     4 a     L  
## 5     5 b     K
```

groups it by y

b. Write down what you think the output will look like, then check if you were correct, and describe what `arrange()` does. Also, comment on how it's different from the `group_by()` in part (a).

```
df |>  
  arrange(y)  
  
## # A tibble: 5 x 3  
##       x y     z  
##   <int> <chr> <chr>  
## 1     1 a     K  
## 2     3 a     L  
## 3     4 a     L  
## 4     2 b     K  
## 5     5 b     K
```

this will sort it alphabetically based on y

c. Write down what you think the output will look like, then check if you were correct, and describe what the pipeline does.

```
df |>  
  group_by(y) |>  
  summarize(mean_x = mean(x))  
  
## # A tibble: 2 x 2  
##   y     mean_x  
##   <chr>   <dbl>  
## 1 a     2.67
```

```
## 2 b      3.5
```

will summarize one row per group

- d. Write down what you think the output will look like, then check if you were correct, and describe what the pipeline does. Then, comment on what the message says.

```
df |> group_by(y, z) |> summarize(mean_x = mean(x))
```

```
## `summarise()` has grouped output by 'y'. You can override using the `.groups`  
## argument.
```

```
## # A tibble: 3 x 3  
## # Groups:   y [2]  
##   y     z   mean_x  
##   <chr> <chr> <dbl>  
## 1 a     K       1  
## 2 a     L     3.5  
## 3 b     K     3.5
```

one row for each combination of y and z

- e. Write down what you think the output will look like, then check if you were correct, and describe what the pipeline does. How is the output different from the one in part (d)?

```
df |>  
  group_by(y, z) |>  
  summarize(mean_x = mean(x), .groups = "drop")
```

```
## # A tibble: 3 x 3  
##   y     z   mean_x  
##   <chr> <chr> <dbl>  
## 1 a     K       1  
## 2 a     L     3.5  
## 3 b     K     3.5
```

dropping groups prevents retention of grouping

- f. Write down what you think the outputs will look like, then check if you were correct, and describe what each pipeline does. How are the outputs of the two pipelines different?

```
df |>  
  group_by(y, z) |>  
  mutate(mean_x = mean(x))
```

```
## # A tibble: 5 x 4  
## # Groups:   y, z [3]  
##       x y     z   mean_x  
##   <int> <chr> <chr> <dbl>  
## 1     1 a     K       1  
## 2     2 b     K     3.5  
## 3     3 a     L     3.5  
## 4     4 a     L     3.5  
## 5     5 b     K     3.5
```

`mutate()` adds a column with repeated summary stats, while `summarize()` condenses groups to single rows